

Course 2 - Gaussian Versus Pareto Distributions

UCLA - Econ 19 - Fall 2018

François Geerolf

Contents

Introduction	1
1 Mathematics of Statistical Distributions	2
1.1 Bell-Shaped Distributions	2
1.2 Pareto Distributions	2
2 Some Real-Life Distributions	2
2.1 Natural Sciences	3
2.2 Cities	5
3 MSAs	8

Introduction

During this course, we shall try to understand a technical passage in Rosen [1983]’s *The American Scholar* piece:

Of particular interest here is an observation, first studied systematically by the great Italian economist Vilfredo Pareto in the late nineteenth century, that the distribution of income contains an unusually large proportion of top earners: that is, among the rich rather than the poor. A visual image will perhaps clarify what is meant by “unusual” in this connection. Imagine a graph plotting IQ scores on the horizontal and the frequency of scores on the vertical. The result is a familiar bell-shaped curve. The peak of the bell occurs at a score arbitrarily scaled at 100 and the curve falls symmetrically on either side of 100. Now picture a similar graph, except with earnings on the horizontal. The resulting curve is unbalanced and nonsymmetrical - a bell that is definitely out of whack. To the left of the modal (peak) value it appears much like the IQ frequency curve. However, to the right of the mode it does not fall as fast as it does to the left. It looks as if someone had stood at the right end of the curve, placed it over his back like a rope, and dragged and stretched it out a very long distance. The upper or right-hand tail of the distribution of income is much thicker than the lower, left-hand tail. The extra weight on the right lends a certain skewness to the distribution of income. What this comes down to is that the distribution of earnings is far from proportionate to the distribution of ability. Amazingly, Pareto’s observations have been qualitatively duplicated in virtually every era of every society for which data on income distributions can be found.

In this passage, Sherwin Rosen draws a sharp distribution between Gaussian distributions on the one hand (characterized by the well known bell-shaped curve) and Pareto distributions on the other hand:

1. “Imagine a graph plotting IQ scores on the horizontal and the frequency of scores on the vertical. The result is a familiar Bell-shaped curve.”
2. “The upper or right-hand tail of the distribution of income is much thicker than the lower, left-hand tail. The extra weight on the right lends a certain skewness to the distribution of income. What this comes down to is that the distribution of earnings is far from proportionate to the distribution of ability.”

We first investigate the mathematics of these different distributions, before proceeding to describing some real-world statistical distributions, and connect them to Bell-shaped curves on the one hand and Pareto distributions on the other hand.

1 Mathematics of Statistical Distributions

In order to understand Sherwin Rosen's above comment, I need to take you through some mathematics. Do not panic ! I am going to take you through everything, and a prerequisite of mathematics from high school should be sufficient.

1.1 Bell-Shaped Distributions

The Bell shape curve is defined by a density function given by:

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-\mu)^2}{\sigma^2}\right).$$

One implication is that the density of a Bell Shaped curve goes very rapidly to zero as x goes to infinity. When premultiplied by any power function x^a , no matter how large a , the density of a Bell-shaped curve still converges to zero, which means that the density is negligible compared to any power function when x to infinity:

$$\text{for all } a > 0, \quad \lim_{x \rightarrow +\infty} x^a f(x) = 0.$$

Intuitively, this means that the Gaussian Distribution goes "very fast" to zero as x becomes large, faster in fact than usual functions which are thought to go very fast to zero (thing, for example of x^{10000} when x becomes large).

Here is a link to the Google Sheets that we created in order to look at the Gaussian distribution. In particular, we were able to plot the density function of a Normal Distribution with $\mu = 0$ and $\sigma = 1$, using the formula above. Note: this Google Sheet is read only. However, you may copy and paste from this Google Sheet, and choose your own values for μ and σ .

1.2 Pareto Distributions

A key feature of the Pareto distribution is that the density distribution does not go as fast to 0 as with the Gaussian Distribution, as x becomes large.

In the context For concreteness, if x is population, then this would mean that there are relatively many cities with a large size, especially when assessed against the average city size, as well as its standard deviation. Similarly, there are relatively many incomes that are much larger than the mean. The Pareto Distribution is in fact defined by:

$$f(x) = a \frac{x_m^a}{x^{a+1}}.$$

For the cumulative distribution function, this implies:

$$1 - F(x) = \left(\frac{x_m}{x}\right)^a.$$

2 Some Real-Life Distributions




Figure 1: BELL SHAPED CURVE

```
pklist <- c("tidyverse", "rvest", "scales", "knitr")
source("https://fgeerolf.github.io/code/load-packages.R")
options(tibble.print_max = 100)
```

2.1 Natural Sciences

Many distributions in the natural sciences are well described by a Bell-shaped curve. In order to illustrate this, let us use the National Longitudinal Surveys (NLS) from the Bureau of Labor Statistics which tracks the income, education, and life circumstances of a large cohort of Americans across several decades. We use `summary` in order to summarise our dataset.

```
height <- read.csv("https://raw.githubusercontent.com/hadley/r4ds/master/data/heights.csv")
height %>%
  summary
```

```
##      earn      height      sex      ed
##  Min.   : 200    Min.   :57.50 female:687  Min.   : 3.0
## 1st Qu.:10000   1st Qu.:64.01  male :505  1st Qu.:12.0
## Median :20000   Median :66.45           Median :13.0
## Mean   :23155   Mean   :66.92           Mean   :13.5
## 3rd Qu.:30000   3rd Qu.:69.85           3rd Qu.:16.0
## Max.   :200000   Max.   :77.05           Max.   :18.0
##      age      race
##  Min.   :18.00 black  :112
## 1st Qu.:29.00 hispanic: 66
## Median :38.00 other   : 25
## Mean   :41.38 white   :989
## 3rd Qu.:51.00
```

```
## Max. :91.00
```


The variable names are pretty self-explanatory. We are in particular interested by the distribution of height, possibly by gender.

```
heights <- height %>%
  select(height) %>%
  unlist %>%
  unname


heights.male <- height %>%
  filter(sex == "male") %>%
  select(height) %>%
  unlist %>%
  unname

heights.female <- height %>%
  filter(sex == "female") %>%
  select(height) %>%
  unlist %>%
  unname
```


```
ggplot() +
  aes(heights) +
  geom_histogram(binwidth = 1, colour = "black", fill = "white") +
  theme_bw()
```



```
ggplot() +
  aes(heights.male) +
  geom_histogram(binwidth = 1, colour = "black", fill = "white") +
  theme_bw()
```



```
ggplot() +
  aes(heights.female) +
  geom_histogram(binwidth = 1, colour = "black", fill = "white") +
  theme_bw()
```



2.2 Cities

During the class, we have used this Google Spreadsheet in order to plot the city size distribution of cities. We note that the result is something that is close to a linear relationship, when the log rank is plotted against the log size, which shows that the distribution is close to Pareto.

We can also download everything in R directly. The data comes from the following Wikipedia entry: List of United States cities by population.

```
data <- "https://en.wikipedia.org/wiki/List_of_United_States_cities_by_population" %>%
  read_html %>%
  html_table(header = TRUE, fill = TRUE)
```

City Size Distribution: Size-Rank Log-Log Plot




Figure 2: CITY SIZE DISTRIBUTION

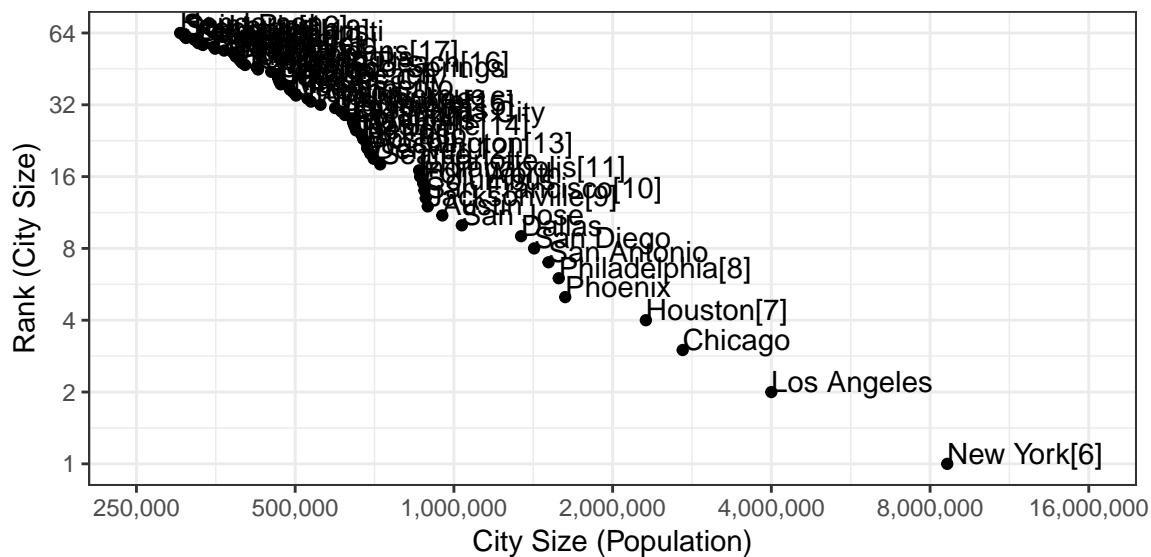
Biggest cities:

```
data[[5]][,c(1:4)] %>%
  as.tibble %>%
  select(rank = "2017rank", "City", state = "State[5]", pop = "2017estimate") %>%
  head(28) %>%
  kable(align = "c")
```

rank	City	state	pop
1	New York[6]	New York	8,622,698
2	Los Angeles	California	3,999,759
3	Chicago	Illinois	2,716,450
4	Houston[7]	Texas	2,312,717
5	Phoenix	Arizona	1,626,078
6	Philadelphia[8]	Pennsylvania	1,580,863
7	San Antonio	Texas	1,511,946
8	San Diego	California	1,419,516
9	Dallas	Texas	1,341,075
10	San Jose	California	1,035,317
11	Austin	Texas	950,715
12	Jacksonville[9]	Florida	892,062
13	San Francisco[10]	California	884,363
14	Columbus	Ohio	879,170
15	Fort Worth	Texas	874,168
16	Indianapolis[11]	Indiana	863,002
17	Charlotte	North Carolina	859,035
18	Seattle	Washington	724,745
19	Denver[12]	Colorado	704,621
20	Washington[13]	District of Columbia	693,972
21	Boston	Massachusetts	685,094

rank	City	state	pop
22	El Paso	Texas	683,577
23	Detroit	Michigan	673,104
24	Nashville[14]	Tennessee	667,560
25	Memphis	Tennessee	652,236
26	Portland	Oregon	647,805
27	Oklahoma City	Oklahoma	643,648
28	Las Vegas	Nevada	641,676

```
data[[5]][,c(1:4)] %>%
  as.tibble %>%
  select(rank = "2017rank", pop = "2017estimate", "City") %>%
  mutate(pop = pop %>% gsub(",", "", .) %>% as.numeric) %>%
  arrange(-pop) %>%
  mutate(rank = 1:n()) %>%
  ggplot(aes(x = pop, y = rank)) + geom_point() + theme_bw() +
  geom_text(aes(label = City), hjust = 0, vjust = 0) +
  scale_y_log10(breaks = 2^(seq(0, 10, 1)),
    limits = c(1, 64)) +
  scale_x_log10(breaks = 250000*2^seq(0, 10, 1),
    limits = c(250000, 16000000),
    labels = comma) +
  ylab("Rank (City Size)") + xlab("City Size (Population)")
```

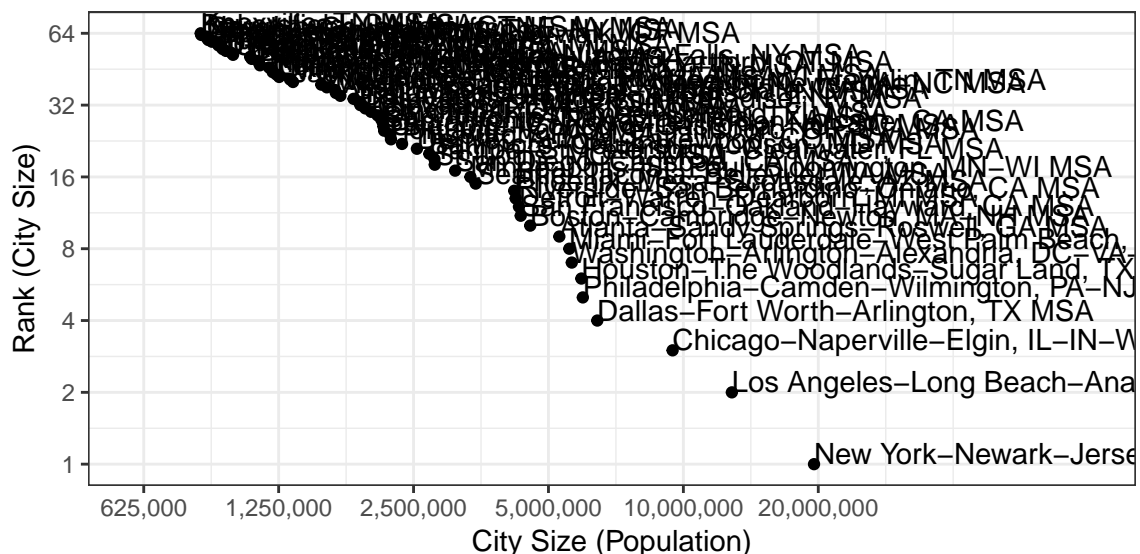


```
data[[5]][,c(1:4)] %>%
  as.tibble %>%
  select(rank = "2017rank", pop = "2017estimate", "City") %>%
  mutate(pop = pop %>% gsub(",", "", .) %>% as.numeric) %>%
  arrange(-pop) %>%
  mutate(rank = 1:n()) %>%
  filter(pop >= 500000) %>%
  ggplot(aes(x = pop, y = rank)) + geom_point() + theme_bw() +
  geom_text(aes(label = City), hjust = 0, vjust = 0) +
  scale_y_log10(breaks = 2^(seq(0, 10, 1)),
    limits = c(1, 28)) +
```



Metropolitan statistical area	2010 Census
Riverside-San Bernardino-Ontario, CA MSA	4,224,851
Detroit-Warren-Dearborn, MI MSA	4,296,250
Seattle-Tacoma-Bellevue, WA MSA	3,439,809
Minneapolis-St. Paul-Bloomington, MN-WI MSA	3,348,859
San Diego-Carlsbad, CA MSA	3,095,313
Tampa-St. Petersburg-Clearwater, FL MSA	2,783,243
Denver-Aurora-Lakewood, CO MSA	2,543,482
Baltimore-Columbia-Towson, MD MSA	2,710,489
St. Louis, MO-IL MSA	2,787,701
Charlotte-Concord-Gastonia, NC-SC MSA	2,217,012
Orlando-Kissimmee-Sanford, FL MSA	2,134,411
San Antonio-New Braunfels, TX MSA	2,142,508
Portland-Vancouver-Hillsboro, OR-WA MSA	2,226,009
Pittsburgh, PA MSA	2,356,285
Sacramento-Roseville-Arden-Arcade, CA MSA	2,149,127
Las Vegas-Henderson-Paradise, NV MSA	1,951,269

```
temp <- data[[2]][,c(2, 4)]
names(temp) <- c("MSA", "pop")

temp %>%
  as.tibble %>%
  mutate(pop = pop %>% gsub(",", "", .) %>% as.numeric) %>%
  arrange(-pop) %>%
  mutate(rank = 1:n()) %>%
  ggplot(aes(x = pop, y = rank)) + geom_point() + theme_bw() +
  geom_text(aes(label = MSA), hjust = 0, vjust = 0) +
  scale_y_log10(breaks = 2^(seq(0, 10, 1)),
               limits = c(1, 64)) +
  scale_x_log10(breaks = 2500000*2^seq(-3, 3, 1),
               limits = c(600000, 80000000),
               labels = comma) +
  ylab("Rank (City Size)") + xlab("City Size (Population)")
```



```
temp %>%
  as.tibble %>%
  mutate(pop = pop %>% gsub(",", "", .) %>% as.numeric) %>%
  arrange(-pop) %>%
  mutate(rank = 1:n()) %>%
  filter(pop >= 500000) %>%
  ggplot(aes(x = pop, y = rank)) + geom_point() + theme_bw() +
  geom_text(aes(label = MSA), hjust = 0, vjust = 0) +
  scale_y_log10(breaks = 2^(seq(0, 10, 1)),
               limits = c(1, 28)) +
  scale_x_log10(breaks = 2500000*2^seq(0, 3, 1),
               limits = c(1900000, 80000000),
               labels = comma) +
  ylab("Rank (City Size)") + xlab("City Size (Population)")
```



References

Sherwin Rosen. The Economics of Superstars. *The American Scholar*, 52(4):449–460, 1983. ISSN 0003-0937.
URL <http://www.jstor.org/stable/41210977>.